



Task#1 (Recursive Query)

CLO 2.6

Q 1: Consider a table named Directory with the following columns:

ID (integer, primary key): Unique identifier for each directory entry.

Name (text): Name of the directory entry.

ParentID (integer, foreign key): ID of the parent directory. Null if the entry is at the root level.

Write a SQL query to retrieve the full path of each directory entry in the Directory table. The output should include the ID, Name, and FullPath columns, where FullPath represents the complete path of each directory entry, starting from the root directory and separated by slashes ("/").

[3 Marks]

The code:

```
1 CREATE TABLE Directory (  
2   ID INT PRIMARY KEY,  
3   Name VARCHAR(100),  
4   ParentID INT,  
5   FOREIGN KEY (ParentID) REFERENCES Directory(ID)  
6 );-- Insert data into the Directory table  
7  
8 INSERT INTO Directory (ID, Name, ParentID) VALUES  
9 (1, 'Root', NULL),  
10 (2, 'F1', 1),  
11 (3, 'F2', 1),  
12 (4, 'SubF1', 2),  
13 (5, 'SubF2', 2),  
14 (6, 'Subsub1', 4),  
15 (7, 'F1.txt', 6),  
16 (8, 'F2.txt', 6),  
17 (9, 'F3.txt', 5);-- Query to retrieve the full path of each directory entry  
18  
19 WITH RECURSIVE DirectoryPath AS (  
20   SELECT ID, Name, Name AS FullPath, ParentID  
21   FROM Directory  
22   WHERE ParentID IS NULL-- Root directory  
23   UNION ALL  
24   SELECT d.ID, d.Name, CONCAT(dp.FullPath, '/', d.Name) AS FullPath, d.ParentID  
25   FROM Directory d  
26   INNER JOIN DirectoryPath dp ON d.ParentID = dp.ID  
27 )  
28 SELECT ID, Name, FullPath  
29 FROM DirectoryPath;
```



The execution:

1:

The screenshot shows the MySQL Workbench interface. The SQL Editor contains the following script:

```
1 CREATE TABLE Directory (  
2   ID INT PRIMARY KEY,  
3   Name VARCHAR(100),  
4   ParentID INT,  
5   FOREIGN KEY (ParentID) REFERENCES Directory(ID)  
6 )  
7 -- Insert data into the Directory table  
8 INSERT INTO Directory (ID, Name, ParentID) VALUES  
9   (1, 'Root', NULL),  
10  (2, 'F1', 1),  
11  (3, 'F2', 1),  
12  (4, 'SubF1', 2),  
13  (5, 'SubF2', 2),
```

The Output tab shows the execution results:

#	Time	Action	Message	Duration / Fetch
1	22:33:36	CREATE TABLE Directory (ID INT PRIMARY KEY, Name VARCHAR(100), ParentID INT, FOREIGN KEY (ParentID) REFERENCES Directory(ID))	0 row(s) affected	0.031 sec

2:

The screenshot shows the MySQL Workbench interface. The SQL Editor contains the following script:

```
8 INSERT INTO Directory (ID, Name, ParentID) VALUES  
9   (1, 'Root', NULL),  
10  (2, 'F1', 1),  
11  (3, 'F2', 1),  
12  (4, 'SubF1', 2),  
13  (5, 'SubF2', 2),  
14  (6, 'Subsub1', 4),  
15  (7, 'F1.txt', 6),  
16  (8, 'F2.txt', 6),  
17  (9, 'F3.txt', 5) -- Query to retrieve the full path of each directory entry  
18 WITH RECURSIVE DirectoryPath AS (  
19   SELECT ID, Name, Name AS FullPath, ParentID  
20   FROM Directory
```

The Output tab shows the execution results:

#	Time	Action	Message	Duration / Fetch
1	22:33:36	CREATE TABLE Directory (ID INT PRIMARY KEY, Name VARCHAR(100), ParentID INT, FOREIGN KEY (ParentID) REFERENCES Directory(ID))	0 row(s) affected	0.031 sec
2	22:34:43	INSERT INTO Directory (ID, Name, ParentID) VALUES (1, 'Root', NULL), (2, 'F1', 1), (3, 'F2', 1), (4, 'SubF1', 2), (5, 'SubF2', 2), (6, 'Subsub1', 4), (7, 'F1.txt', 6), (8, 'F2.txt', 6), (9, 'F3.txt', 5)	9 row(s) affected Records: 9 Duplicates: 0 Warnings: 0	0.015 sec



Output:

MySQL Workbench

Local instance MySQL80 x

File Edit View Query Database Server Tools Scripting Help

Navigator

reemarenad* SQL File 2* SQL File 4*

Limit to 1000 rows

SQL

22 UNION ALL
23 SELECT d.ID, d.Name, CONCAT(dp.FullPath, '/', d.Name) AS FullPath, d.ParentID
24 FROM Directory d
25 INNER JOIN DirectoryPath dp ON d.ParentID = dp.ID
26)
27 SELECT ID, Name, FullPath
28 FROM DirectoryPath;

Result Grid

ID	Name	FullPath
1	Root	Root
2	F1	Root/F1
3	F2	Root/F2
4	SubF1	Root/F1/SubF1
5	SubF2	Root/F1/SubF1/SubF2
6	Subsub1	Root/F1/SubF1/Subsub1
9	F3.txt	Root/F1/SubF2/F3.txt
7	F1.txt	Root/F1/SubF1/Subsub1/F1.txt

Result 3 x

Read Only Context Help Snippets

Output

Action Output

#	Time	Action	Message	Duration / Fetch
1	22:33:36	CREATE TABLE Directory (ID INT PRIMARY KEY, Name VARCHAR(100), Parent...	0 row(s) affected	0.031 sec
2	22:34:43	INSERT INTO Directory (ID, Name, ParentID) VALUES (1, 'Root', NULL), (2, 'F1', 1)...	9 row(s) affected Records: 9 Duplicates: 0 Warnings: 0	0.015 sec
3	22:35:27	WITH RECURSIVE DirectoryPath AS (SELECT ID, Name, Name AS FullPath, Pare...	9 row(s) returned	0.000 sec / 0.000 sec

Object Info Session

Automatic context help is disabled. Use the toolbar to manually get help for the current caret position or to toggle automatic help.

Task#2 (B+ Tree Indexing)

CLO 2.2

Q.2: Consider the B+ tree shown in Figure 1, do the following:





Figure 1: A B+ Tree


a) Add key 195, and update the B+ tree.

[2 Marks]



b) Delete key 200 from the tree obtained from Q1 (a)

[3 Marks]



Task#3 (Transactions and Schedules in Databases)

CLO 2.4

Q.3(a): You are given the following schedules involving three transactions (T1, T2, T3) performing operations on database items (A, B, C, D):

[2 Marks]

**Schedule S1:**


T1: R(A), R(B), W(C)
 T2: W(A), R(B), R(C)
 T3: W(A), W(B), R(D)

T1 → T2: Conflict 5

T1 → T3: Conflict R(A) in T1, W(A) T3

T2 → T3: Conflict 7

T1	T2	T3
R(A)	W(A)	W(A)
R(B)	R(B)	W(B)
W(C)	R(C)	R(D)



T1 R(A) → T2 W(A)

T1 R(A) → T3 W(A)

T1 R(B) → T3 W(B)

T1 W(C) → T2 R(C)

T2 W(A) → T3 W(A)


T2 R(B) → T3 W(B)

S1: Conflict-serializable because there is no cycle.

Schedule S2:

T1: R(A), R(C), W(C), R(D)
 T2: R(A), W(B), R(C), W(D)
 T3: W(B), W(A), W(D), R(C)

T1	T2	T3
R(A)	R(A)	W(B)
R(C)	W(B)	W(A)
W(C)	R(C)	W(D)
R(D)	W(D)	R(C)



T1 R(A) → T3 W(A)

T1 W(C) → T2 R(C)

T1 R(D) → T2 W(D)

T2 R(A) → T3 W(A)

T3 W(B) → T2 W(B)

T3 W(D) → T2 W(D)


T3 W(D) → T1 R(D)

S2: Not conflict-serializable because there is a cycle.

**Schedule S3:**

T1: R(A), R(B), R(C), W(D)
 T2: W(A), W(B), W(C), R(D)
 T3: W(A), R(B), W(C), W(D)

T1	T2	T3
R(A)	W(A)	W(A)
R(B)	W(B)	R(B)
R(C)	W(C)	W(C)
W(D)	R(D)	W(D)



T1 R(A) → T2 W(A)
 T1 R(A) → T3 W(A)
 T1 R(B) → T2 W(B)
 T1 R(C) → T2 W(C)
 T1 R(C) → T3 W(C)
 T1 W(D) → T2 R(D)
 T1 W(D) → T3 W(D)

S3: Conflict-serializable because there is no cycle.


Your task is to verify the conflict serializability of each schedule and determine whether they are conflict-serializable or not. If it's not, explain the conflicting operations that violate conflict-serializability.

Q.3(b): Consider the following two schedules; Schedule A and Schedule B, as given below: **[2 Marks]**

Schedule A:

T1: R1(X); W2(Y); W1(Y); R3(Z); R1(Y); W2(X); W2(Z); W3(X)
 T2: R1(Y); W1(X); R2(Y); R2(X)
 T3: W1(Z); W3(Y); R3(X); R2(Z); R1(Z)

T1	T2	T3
R(X)		
	W(Y)	
W(Y)		
		R(Z)
R(Y)		
	W(X)	
	W(Z)	
		W(X)
R(Y)		
W(X)		
	R(Y)	
	R(X)	
W(Z)		
		W(Y)
		R(X)
	R(Z)	
R(Z)		



Schedule "A" is not conflict-serializable.




Schedule B:

T1: R1(X); R2(Y); W1(X); R2(Z); W2(X); W3(Z); R3(X)

T2: R1(Y); R3(Z); W1(Y); R2(X); R2(Z)

T3: R1(Z); W2(Y); R3(X); W1(Z)

T1	T2	T3
R(X)		
	R(Y)	
W(X)		
	R(Z)	
	W(X)	
		W(Z)
		R(X)
R(Y)		
		R(Z)
W(Y)		
	R(X)	
	R(Z)	
R(Z)		
	W(Y)	
		R(X)
W(Z)		



Schedule "B" is not conflict-serializable.

Draw the precedence graph for Schedule A and Schedule B. State whether Schedules are conflict-serializable or not, and explain your reasoning.

Q.3(c): Consider the following two transactions, T1 and T2, involved in a multi-user database environment:

Transaction#1

```
read(X);
read(Y);
if X = 0 then Y := Y + 1;
write(Y).
```




Answer Transaction#1:

lock-S(X)

read (X)

lock-X(Y)

read (Y)

if $X = 0$ then $Y := Y + 1$

write(Y)

unlock(X)

unlock(Y)

Transaction#2

read(Y);

read(X);

if $Y = 0$ then $X := X + 1$;

write(X).

Answer Transaction#2:

lock-S(Y)

read(Y)

lock-X(X)

read (X)

if $Y = 0$ then $X := X+1$

write(X)

unlock(Y)

unlock(X)

Modify transactions T1 and T2 by adding appropriate lock and unlock instructions to ensure they adhere to the two-phase locking protocol.